

Model Optimization and Tuning Phase Template

Date	10 July 2024
Team ID	SWTID1720075414
Project Title	Panic Disorder Detection
Maximum Marks	10 Marks

Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase involves refining machine learning models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

Hyperparameter Tuning Documentation (6 Marks): (Not mandatory for this project)

Model	Tuned Hyperparameters	Optimal Values
Decision Tree	<pre> param_grid = { 'criterion': ['gini', 'entropy'], 'max_depth': [None, 5,10,15], 'min_samples_split': [2,3,10], 'min_samples_leaf': [1,2,3], 'max_features': [None, 'sqrt', 'log2'] } #creating a decision tree classifier dt_classifier = DecisionTreeClassifier(random_state=1234) #Create GridSearchCV object grid_search = GridSearchCV(dt_classifier, param_grid, cv=5, verbose = 1, n_jobs=-1) #Fit the data to perform grid search grid_search.fit(x_res_train[fts], y_res_train) </pre>	<pre> #Print the best hyperparameters print("Best Hyperparameters:", grid_search.best_params_) print("Best Score", grid_search.best_score_) </pre> <p>Fitting 5 folds for each of 216 candidates, totalling 1080 Best Hyperparameters: {'criterion': 'gini', 'max_depth': N Best Score 0.9914015567048008</p>
Random Forest	<pre> param_grid = { 'n_estimators': [50,100,200], 'max_depth': [None, 5,10], 'min_samples_split': [2,5,10], 'min_samples_leaf': [1,2,4], 'max_features': ['sqrt', 'log2'] } #creating a Random Forest Classifier rf_classifier = RandomForestClassifier(random_state=1234) #Create GridSearchCV object grid_search = GridSearchCV(rf_classifier, param_grid=param_grid, cv=5, verbose = 1, n_jobs=4) #Fit the data to perform grid search grid_search.fit(x_res_train[fts], y_res_train) </pre>	<pre> #Print the best hyperparameters print("Best Hyperparameters:", grid_search.best_params_) print("Best Score", grid_search.best_score_) </pre> <p>Fitting 5 folds for each of 162 candidates, totalling 810 Best Hyperparameters: {'max_depth': None, 'max_features': Best Score 0.9909784255341378</p>

XGBoost	<pre> param_grid = { 'min_child_weight': [10,20], 'max_depth': [0,1,5,2,0], 'gamma': [0.6,0.8,0.9], 'max_depth': [4,5,6], } xgb = xgboost.XGBClassifier(learning_rate=0.5, n_estimators=100, objective='binary:logistic', nthread=3) fitmodel = GridSearchCV(xgb, param_grid=param_grid, cv=5, refit=True, scoring = 'accuracy', n_jobs=4, verbose=3) #Fit the data to perform grid search fitmodel.fit(x_res_train[fts], y_res_train) </pre>	<pre> #Print the best hyperparameters print("Best Hyperparameters:", fitmodel.best_params_) print("Best Score", fitmodel.best_score_) Fitting 5 folds for each of 54 candidates, totalling 270 fits Best Hyperparameters: {'colsample_bytree': 0.9, 'gamma': 0, 'm Best Score 0.9899526187118007 </pre>
KNN	<pre> param_grid = { 'n_neighbors': [3,5,7], 'weights': ['uniform', 'distance'], 'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'], 'leaf_size': [20,30,40], 'p': [1,2] } knn = KNeighborsClassifier() #Create GridSearchCV object grid_search = GridSearchCV(estimator=knn, param_grid=param_grid, cv=5, verbose = 1, n_jobs=4) #Fit the data to perform grid search grid_search.fit(x_res_train[fts], y_res_train) </pre>	<pre> # Print the best parameters and best score print("Best Parameters: ", grid_search.best_params_) print("Best Score: ", grid_search.best_score_) Fitting 5 folds for each of 144 candidates, totalling Best Parameters: {'algorithm': 'brute', 'leaf_size': Best Score: 0.7771295215869312 </pre>

Performance Metrics Comparison Report (2 Marks):

Model	Optimized Metric															
Decision Tree	<pre>y_pred = grid_search.best_estimator_.predict(x_test[fts]) print(confusion_matrix(y_test,y_pred)) print(classification_report(y_test,y_pred)) print("SCORE:",grid_search.best_estimator_.score(x_test[fts],y_test))</pre>															
	<pre>[[18898 261] [4 837]]</pre>															
	<table><tr><td></td><td>precision</td><td>recall</td><td>f1-score</td><td>support</td></tr><tr><td>0</td><td>1.00</td><td>0.99</td><td>0.99</td><td>19159</td></tr><tr><td>1</td><td>0.76</td><td>1.00</td><td>0.86</td><td>841</td></tr></table>		precision	recall	f1-score	support	0	1.00	0.99	0.99	19159	1	0.76	1.00	0.86	841
		precision	recall	f1-score	support											
	0	1.00	0.99	0.99	19159											
	1	0.76	1.00	0.86	841											
	<table><tr><td>accuracy</td><td></td><td></td><td>0.99</td><td>20000</td></tr><tr><td>macro avg</td><td>0.88</td><td>0.99</td><td>0.93</td><td>20000</td></tr><tr><td>weighted avg</td><td>0.99</td><td>0.99</td><td>0.99</td><td>20000</td></tr></table>	accuracy			0.99	20000	macro avg	0.88	0.99	0.93	20000	weighted avg	0.99	0.99	0.99	20000
	accuracy			0.99	20000											
	macro avg	0.88	0.99	0.93	20000											
	weighted avg	0.99	0.99	0.99	20000											
SCORE: 0.98675																

Random Forest	<pre>y_pred = grid_search.best_estimator_.predict(x_test[fts]) print(confusion_matrix(y_test,y_pred)) print(classification_report(y_test,y_pred)) print("SCORE:",grid_search.best_estimator_.score(x_test[fts],y_test))</pre> <pre>[[18848 311] [14 827]]</pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>0</td><td>1.00</td><td>0.98</td><td>0.99</td><td>19159</td></tr><tr><td>1</td><td>0.73</td><td>0.98</td><td>0.84</td><td>841</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.98</td><td>20000</td></tr><tr><td>macro avg</td><td>0.86</td><td>0.98</td><td>0.91</td><td>20000</td></tr><tr><td>weighted avg</td><td>0.99</td><td>0.98</td><td>0.98</td><td>20000</td></tr></table> <p>SCORE: 0.98375</p>		precision	recall	f1-score	support	0	1.00	0.98	0.99	19159	1	0.73	0.98	0.84	841	accuracy			0.98	20000	macro avg	0.86	0.98	0.91	20000	weighted avg	0.99	0.98	0.98	20000
	precision	recall	f1-score	support																											
0	1.00	0.98	0.99	19159																											
1	0.73	0.98	0.84	841																											
accuracy			0.98	20000																											
macro avg	0.86	0.98	0.91	20000																											
weighted avg	0.99	0.98	0.98	20000																											
KNN	<pre>y_pred=grid_search.best_estimator_.predict(x_test[fts]) print(confusion_matrix(y_test,y_pred)) print(classification_report(y_test,y_pred)) print("SCORE:",grid_search.best_estimator_.score(x_test[fts],y_test))</pre> <pre>[[13375 5784] [565 276]]</pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>0</td><td>0.96</td><td>0.70</td><td>0.81</td><td>19159</td></tr><tr><td>1</td><td>0.05</td><td>0.33</td><td>0.08</td><td>841</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.68</td><td>20000</td></tr><tr><td>macro avg</td><td>0.50</td><td>0.51</td><td>0.44</td><td>20000</td></tr><tr><td>weighted avg</td><td>0.92</td><td>0.68</td><td>0.78</td><td>20000</td></tr></table> <p>SCORE: 0.68255</p>		precision	recall	f1-score	support	0	0.96	0.70	0.81	19159	1	0.05	0.33	0.08	841	accuracy			0.68	20000	macro avg	0.50	0.51	0.44	20000	weighted avg	0.92	0.68	0.78	20000
	precision	recall	f1-score	support																											
0	0.96	0.70	0.81	19159																											
1	0.05	0.33	0.08	841																											
accuracy			0.68	20000																											
macro avg	0.50	0.51	0.44	20000																											
weighted avg	0.92	0.68	0.78	20000																											
XGBoost	<pre>y_pred = fitmodel.best_estimator_.predict(x_test[fts]) print(confusion_matrix(y_test,y_pred)) print(classification_report(y_test,y_pred)) print("SCORE:",fitmodel.best_estimator_.score(x_test[fts],y_test))</pre> <pre>[[16546 2613] [7 834]]</pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>0</td><td>1.00</td><td>0.86</td><td>0.93</td><td>19159</td></tr><tr><td>1</td><td>0.24</td><td>0.99</td><td>0.39</td><td>841</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.87</td><td>20000</td></tr><tr><td>macro avg</td><td>0.62</td><td>0.93</td><td>0.66</td><td>20000</td></tr><tr><td>weighted avg</td><td>0.97</td><td>0.87</td><td>0.90</td><td>20000</td></tr></table> <p>SCORE: 0.869</p>		precision	recall	f1-score	support	0	1.00	0.86	0.93	19159	1	0.24	0.99	0.39	841	accuracy			0.87	20000	macro avg	0.62	0.93	0.66	20000	weighted avg	0.97	0.87	0.90	20000
	precision	recall	f1-score	support																											
0	1.00	0.86	0.93	19159																											
1	0.24	0.99	0.39	841																											
accuracy			0.87	20000																											
macro avg	0.62	0.93	0.66	20000																											
weighted avg	0.97	0.87	0.90	20000																											

Final Model Selection Justification (2 Marks):

Final Model	Reasoning
Decision Tree	Decision Tree Model was chosen because it displayed an incredible amount of accuracy when run through hyperparameter tuning. It was able to reduce overfitting very well, which fits the project perfectly, making it the perfect for selecting it as the final model for this project.